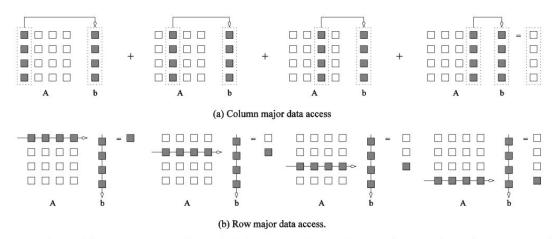
Discussion #3:

WS #5, Lab #1, and Memory Fence

Daniel Tan, Jason Lau, and Jason Cong

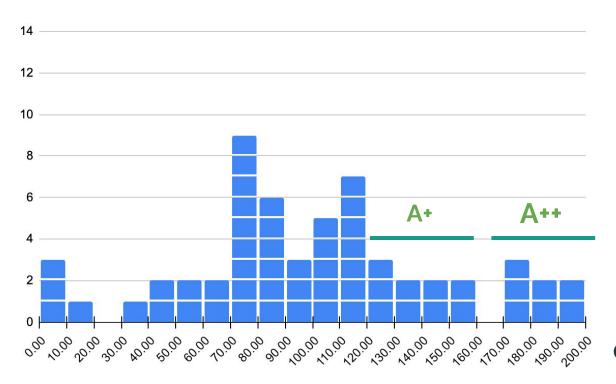


#### Worksheet #5



Multiplying a matrix with a vector: (a) multiplying column-by-column, keeping a running sum; (b) computing each element of the result as a dot product of a row of the matrix with the vector.

#### Lab #1: Performance Distribution



#### A++ Students

1	<u>Poggers</u>	199.43	
2	<u>Teem</u>	192.37	
3	<u>sumedha</u>	182.95	
4	Speed Cola	181.95	
5	<u>m</u>	178.73	
6	<u>AT</u>	176.37	
7	Yingge He	176.14	

#### **GFLOPS**

## Lab #1: Things to Consider

- Which loops to parallelize? Outer loop
- Static, dynamic or guided scheduling?
- How many threads? 8 threads

```
#pragma omp parallel for schedule(static)
for (int i = 0; i < kI; i++)
    for (int j = 0; j < kJ; j++)
        for (int k = 0; k < kK; k++)
            c[i][j] += a[i][k] * b[k][j];</pre>
```

**0.74**GFLOPS

3.89X Sequential

3.89X Improvements

#### Lab #1: Things to Consider

- Which loops to parallelize? Outer loop
- Static, dynamic or guided scheduling?
- How many threads? 8 threads
- Row major access or column major?

**36.74** GFLOPS

193X Sequential

49.6X Improvements

# Lab #1: Why Blocked Design?

1	2	3	4
1	2	3	4
2	4	6	8
3	6	9	12
4	8	12	16
	2	1 2 2 4 3 6	1 2 3 2 4 6 3 6 9

 $Figure: Intel\ https://software.intel.com/content/www/us/en/develop/articles/efficient-use-of-tiling.html. The property of t$ 

# Lab #1: Why Blocked Design?

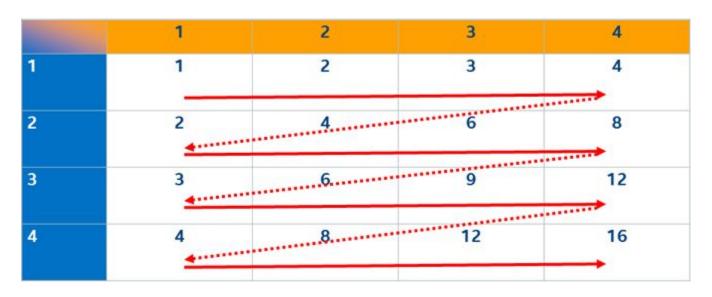


Figure: Intel https://software.intel.com/content/www/us/en/develop/articles/efficient-use-of-tiling.html

# Lab #1: Why Blocked Design?

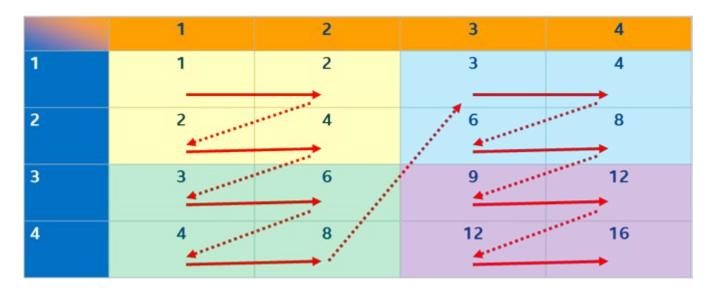
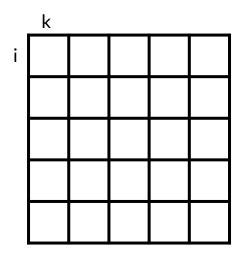
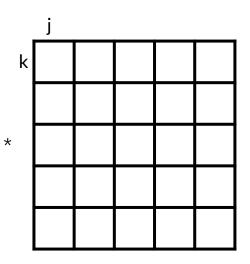


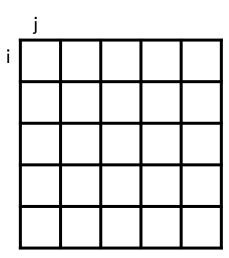
Figure: Intel https://software.intel.com/content/www/us/en/develop/articles/efficient-use-of-tiling.html

#### Lab #1: Blocked Design?

8192 numbers in cache

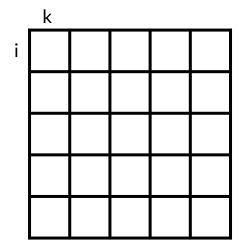


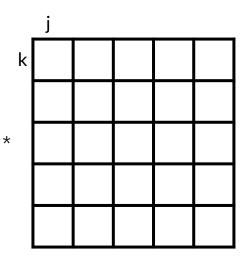


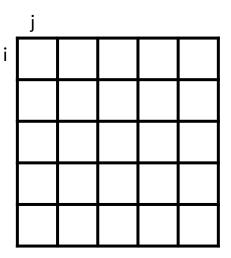


#### Lab #1: Blocked Design?

8192 numbers in cache







#### Lab #1: Things to Consider

Blocked design

```
#pragma omp parallel for schedule(static)
for (int ii = 0; ii < kI; ii += 64)
for (int kk = 0; kk < kK; kk += 32)
for (int jj = 0; jj < kJ; jj += 64)
for (int i = 0; i < kI; i++)
    for (int k = 0; k < kK; k++)
        for (int j = 0; j < kJ; j++)
        c[i][j] += a[i][k] * b[k][j];</pre>
```

### Lab #1: Things to Consider

Blocked design

```
#pragma omp parallel for schedule(static)
for (int ii = 0; ii < kI; ii += 64)
for (int kk = 0; kk < kK; kk += 32)
for (int jj = 0; jj < kJ; jj += 64)
for (int i = ii; i < ii + 64; i++)
    for (int k = kk; k < kk + 32; k++)
        for (int j = jj; j < jj + 64; j++)
        c[i][j] += a[i][k] * b[k][j];</pre>
```

**63.94** GFLOPS

337X Sequential

1.74X Improvements

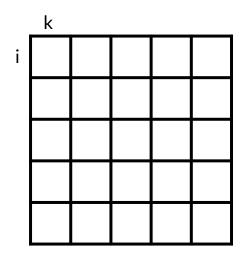
20/25 95/100

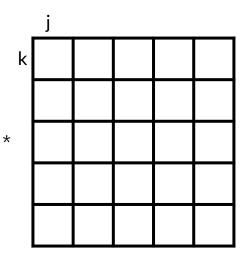
#### Lab #1: Blocked Design?

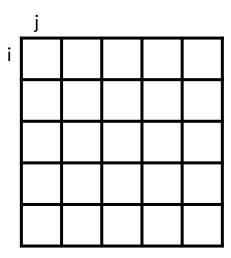
8192 numbers in cache

```
#pragma omp parallel for schedule(static)
for (int i = 0; i < kI; i++)
    for (int k = 0; k < kK; k++)
        for (int j = 0; j < kJ; j++)
            c[i][j] += a[i][k] * b[k][j];</pre>
```

$$B_i = 64$$
  
 $B_j = 1024$   
 $B_k = 8$ 







#### Lab #1: Things to Consider

Blocked design

```
#pragma omp parallel for schedule(static)
for (int ii = 0; ii < kI; ii += 64)
for (int kk = 0; kk < kK; kk += 8)
for (int jj = 0; jj < kJ; jj += 1024)
for (int i = ii; i < ii + 64; i++)
    for (int k = kk; k < kk + 8; k++)
        for (int j = jj; j < jj + 1024; j++)
        c[i][j] += a[i][k] * b[k][j];</pre>
```

**85.09**GFLOPS

448X Sequential

1.33X Improvements

- Blocked design
- Put data in register

- Blocked design
- Put data in register

**104.6** GFLOPS

**547**X Sequential

1.23X Improvements

- Blocked design
- Put data in register

- Blocked design
- Put data in register

- Blocked design
- Put data in register

```
#pragma omp parallel for schedule(static)
for (int ii = 0; ii < kI; ii += 64)
for (int kk = 0; kk < kK; kk += 8)
for (int jj = 0; jj < kJ; jj += 1024)
for (int i = ii; i < ii + 64; i++)
    for (int j = jj; j < jj + 1024; j++) {
        float reg = c[i][j];
        for (int k = kk; k < kk + 8; k++)
            reg += a[i][k] * b[k][j];
        c[i][j] = reg;
}</pre>
```

**144.1** GFLOPS

768X Sequential

1.65X Improvements

 Check all possibilities and fine tune

```
#pragma omp parallel for schedule(static) num_threads(4)
for (int ii = 0; ii < kI; ii += 64)
for (int jj = 0; jj < kJ; jj += 1024)
for (int kk = 0; kk < kK; kk += 8)
for (int i = 0; i < 64; i++)
    for (int j = 0; j < 1024; j++) {
        float reg = c[i + ii][j + jj];
        for (int k = 0; k < 8; k++)
             reg += a[i+ii][k+kk] * b[k+kk][j+jj];
        C[i + ii][j + jj] = reg;
    }</pre>
```

**204.3** GFLOPS

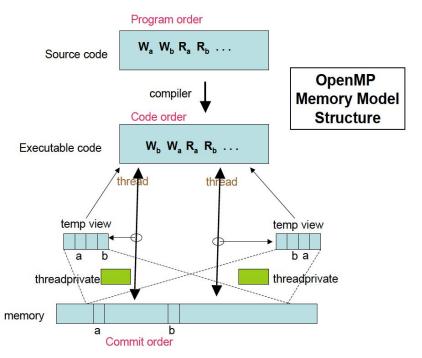
1075X Sequential

1.42X Improvements

## **Memory Fence in OpenMP**

- Relaxed-consistency, shared-memory model
  - All threads have access to the *memory* store and retrieve variables
  - Each thread is allowed to have its own temporary view of the memory
    - If multiple threads write to the same shared variable without synchronization, the resulting value of the variable in memory is unspecified
    - If at least one thread reads from a shared variable and at least one thread writes to it without synchronization, the value seen by any reading thread is unspecified
  - Threadprivate memory must not be accessed by other threads

## **OpenMP Memory Model Structure**



Hoeflinger, de Supinski, "The OpenMP Memory Model," IWOMP 2005.

### **Memory Fences in OpenMP**

- #pragma omp flush (<variable-list>)
  - oAll the shared variables will be flushed if not given the list
  - Automatically inserted at most points where it is needed
    - barrier
    - parallel upon entry and exit
    - critical upon entry and exit
    - o ordered upon entry and exit
    - o for upon exit
    - sections upon exit
    - single upon exit
    - Read/write an object with a volatile-qualified type
    - Section 2.7.5 in OpenMP Specification 2.5

#### Example: flush Pragma

```
int flag[MAX];
int i:
for (i = 0; i < MAX; i++)
    flag[i] = 0:
#pragma omp parallel shared(flag)
      int ID = omp_get_thread_num();
      do_a_wholebunch(ID);
      flag[ID] = 1:
#pragma omp flush(flag[ID]);
      while (!flag[neighbor(ID)]) {
#pragma omp flush(flag[neighbor(ID)]);
      do_more_stuff();
```

- How thread A reads the shared variable x written by thread B?
  - Thread B writes the value to the variable x
  - Thread B flushes the variable x
  - Thread A flushes the variable x
  - ∘Thread A read the value from the variable *x*

#### Importance of Variable List in Flush

```
Incorrect example:
                             a = b = 0
     thread 1
                                               thread 2
      b = 1
                                                a = 1
      flush (b)
                                               flush (a)
      flush (a)
                                               flush (b)
      if (a == 0) then
                                                if (b == 0) then
         critical section
                                                   critical section
      end if
                                                end if
```

```
Correct example:

a = b = 0

thread 1

thread 2

b = 1

flush(a,b)

if (a == 0) then

critical section
end if

thread 2

a = 1

flush(a,b)

if (b == 0) then

critical section
end if
```

The flush operation has two effects:

- I. make the temporary view and share memory consistent, and
- II. impose some ordering constraints on the memory accesses.
  In the first implementation, flush(b) in thread 1 and flush(a) in thread 2 might be moved to the end by the compiler, resulting both of them entering the critical sections at the same time.

Q&A